Unsupervised Learning of Cross-lingual Word Embedding and Its Application to Machine Translation

Jiahui Geng

jiahui.geng@rwth-aachen.de

Master Thesis Mid-term Talk
August 6, 2018

Human Language Technology and Pattern Recognition
Lehrstuhl für Informatik 6
Computer Science Department
RWTH Aachen University, Germany

Outline

Introduction

Literature

Cross-lingual word embedding

- Supervised learning
- Unsupervised learning

Sentence Translation with cross-lingual word embedding

- Context-aware beam search
- Denoising autoencoder

Experiments

Outlook

Introduction

Motivation

- Building a machine translation system requires lots of bilingual data
- Cross-lingual word embedding offers elegant word matches between languages
- Unsupervised MT relies on back-translation which needs a long training time

Goals

- Study training details of cross-lingual word embedding
- Build a good unsupervised MT efficiently: combine with other models
- Improve the unsupervised learning algorithm for cross lingual word embedding

Literature

Unsupervised cross-lingual embedding

- [Conneau & Lample^{+} 17] Word translation without parallel data
  - Implementation of GANs: discriminator trained to distinguish between two distributions while generator fools discriminator
- [Artetxe & Labaka^{+} 17a] Learning bilingual word embeddings with (almost) no bilingual data
  - A self-learning framework combining embedding mapping and dictionary induction techniques, needs seed dictionary to start
- [Hoshen & Wolf 18] An Iterative Closest Point Method for Unsupervised Word Translation
  - Iterative closest point method for embedding mapping learning

Literature

Unsupervised machine translation

- [Artetxe & Labaka^{+} 17b] Unsupervised Neural Machine Translation
- [Lample & Denoyer^{+} 17] Unsupervised Machine Translation Using Monolingual Corpora Only
  - Seq2seq model with shared encoder and decoder for both languages, also with denoising autoencoder and back-translation
- [Artetxe & Labaka^{+} 17b] Phrase-Based & Neural Unsupervised Machine Translation
  - Simplifies the architecture and loss function for unsupervised NMT and propose a phrase-based SMT with back-translation

Cross-lingual Word Embedding

Definition

- Word embedding of multiple languages in a joint embedding space
- Linear mapping from source embedding to target embedding (this work)

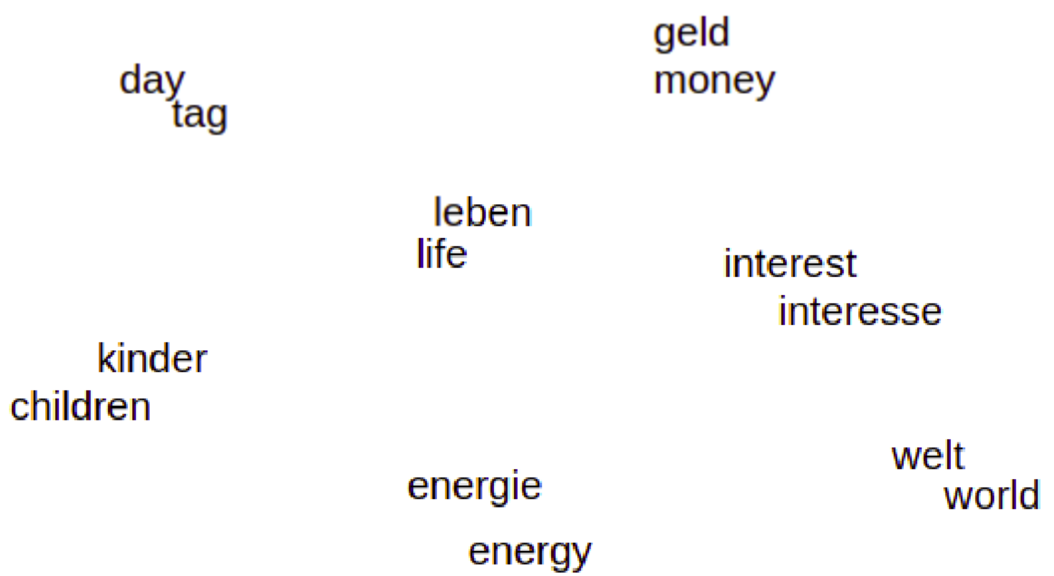

cross-lingual word embedding

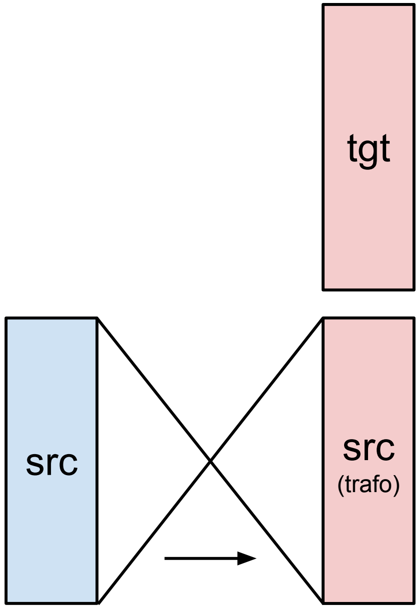

cross-lingual mapping

Cross-lingual Word Embedding

Roles in unsupervised neural machine translation

- Shared latent representations
  - Shared encoder for producing a language independent representation
- As word or phrase table for translation

This work

- Formulate a straightforward way to combine a language model with cross-lingual word similarities

Training Methods

- Mapping-based approaches (this work)
- Pseudo-multi-lingual corpora-based approaches
- Joint methods

Cross-lingual Word Embedding

Mapping based approaches

- Learn monolingual embedding separately
  - Skip-gram model
- Learn linear mapping between embedding spaces
  - Supervised learning
    - Procrustes analysis
  - Semi-supervised learning
    - Iterative self-learning framework
  - Unsupervised learning
    - Adversarial learning
- Synthetic dictionary induction
  - Nearest neighbor search

Monolingual Embedding

`Fasttext` [Bojanowski & Grave+ 17]

- Essentially an extension of skip-gram/CBOW model
- Treat each word as compound of character n-gram
- Learn the internal structure of words
- Score function between context word c and current word w

$$s(w, c) = \sum_{g \in G_w} z_g^T c$$

  - G_w: set of n-gram appears in w
  - z_g,v_c: the corresponding embedding

Supervised Learning

Assume given

- **Word embedding**
  trained independently for each language on monolingual corpora
- **Bilingual dictionary**
  a known dictionary with pairs of words $\{f, e\}$ size N

Learn a linear mapping $W \in \mathbb{R}^{d \times d}$ such that

$$W^* = \underset{W \in \mathbb{R}^{d \times d}}{\operatorname{argmin}} \sum_{n=1}^{N} \|W f_i - e_i\|^2$$

- d **: Dimension of embedding**
- $f_n, e_n \in \mathbb{R}^d$**: the embedding pair of corresponding word pair in the dictionary**

Procrustes Analysis

Constrain W to be an orthogonal matrix

- Enforce monolingual invariance
- Simplify the problem as the Procrustes problem
  - A closed-form solution obtained from SVD
  - $\boldsymbol{E}$, $\boldsymbol{F} \in \mathbb{R}^{d*N}$ denotes embedding projection of word pairs $\{e, f\}$

$$\boldsymbol{W}^* = \operatorname*{argmin}_{\boldsymbol{W} \in \mathbb{R}^{d \times d}} \|\boldsymbol{W}\boldsymbol{F} - \boldsymbol{E}\|_{\boldsymbol{F}}^2 = \boldsymbol{U}\boldsymbol{V}^{\boldsymbol{T}}$$

$$\boldsymbol{U}\boldsymbol{\Sigma}\boldsymbol{V}^{\boldsymbol{T}} = \mathbf{SVD}(\boldsymbol{E}\boldsymbol{F}^{\boldsymbol{T}})$$

- Can be efficiently computed in linear time w.r.t. seed dictionary size N

Semi-supervised Learning

Problem

- Large dictionary not readily available for many language pairs

Self-learning framework

1. Given source and target embedding $\mathcal{F}$ $\mathcal{E}$, seed dictionary D
2. Learn mapping with dictionary
3. Induce dictionary according to mapping
4. Repeat step 2, 3 until converges

Performance

- Model works with initial dictionary
- Achieve comparable accuracy as supervised method
- Stuck in a poor local optimum without initial dictionary

Unsupervised Learning

Methods

- Learn bilingual embeddings without any bilingual evidence (this work)
  - Adversarial training [Conneau & Lample^{+} 17]
- Design the seed dictionary
  - Shared words, digits and cognates [Artetxe & Labaka^{+} 17a]
  - Design heuristics to build the seed dictionary [Hoshen & Wolf 18] [Artetxe & Labaka^{+} 18]

Adversarial Training

Model

- $\mathcal{F} = \left\{f_1, \ldots f_{V_f}\right\}$ and $\mathcal{E} = \{e_1, \ldots e_{V_e}\}$: set of embeddings, not parallel
- Discriminator is trained to discriminate Wf_n and e_n with f_n, e_n randomly sampled from $\mathcal{F}, \mathcal{E}$
- Generator W is trained to prevent the discriminator from making accurate prediction

Discriminator loss

$$\mathcal{L}_D(\theta_D|W) = -\frac{1}{N}\sum_{n=1}^{N} \log P_{\theta_D}(source = 1|Wf_n) - \frac{1}{M}\sum_{m=1}^{M} \log P_{\theta_D}(source = 0|e_m)$$

Generator loss

$$\mathcal{L}_D(W|\theta_D) = -\frac{1}{N}\sum_{n=1}^{N} \log P_{\theta_D}(source = 0|Wf_n) - \frac{1}{M}\sum_{m=1}^{M} \log P_{\theta_D}(source = 1|e_m)$$

Dictionary Induction

Nearest neighbor search

- Hubness problem: some points (hubs) tends to be nearest neighbors of many points in high-dimensional space

Cross-domain Similarity Local Scaling (CSLS)

- Penalize the similarity score of hubs
  - $N_T(Wf)$: target neighbours for mapped source embedding
  - $r_T(Wf)$: penalty for hubness

$$r_T(Wf) = \frac{1}{K} \sum_{e \in N_T(Wf)} \cos(Wf, e)$$

$$\mathrm{CSLS}(Wf, e) = 2\cos(Wf, e) - r_T(Wf) - r_S(e)$$

Bidirection dictionary induction

- Unidirectional dictionary might lead to local optima
- Include only the mutual nearest neighbors

Sentence Translation

Context-aware Beam Search

- **Language model**

Denoising Autoencoder

- **Insertion**
- **Deletion**
- **Reordering**

Context-aware beam search

Given a history h of target word before e, the score of e to be the translation of f:

$$L(e; f, h) = \lambda_{emb} \log q(f, e) + \lambda_{LM} \log p(e|h)$$

- Lexicon score $q(f, e) \in [0, 1]$ defined as:

$$q(f, e) = \frac{d(f, e) + 1}{2}$$

  where $d(f, e) \in [-1, 1]$ cosine similarity between f and e. In experiments, lexicon score from linear scaling works better than others, e.g. sigmoid or softmax

- Empirically set λ_{emb} as 1, λ_{LM} as 0.1

Denoising

Basic idea

- Model $noise(e_1^I)$ by injecting artificial noise into clean sentences e_1^I
- Neural network learns to restore more smooth sentence from word-by-word translation

Training criterion

$$L = E_{e_1^I \in E}[-log(e_1^I | noise(e_1^I))]$$

- E denotes target corpus.
- In Seq2Seq training, e_1^I as label, $noise(e_1^I)$ as input
- Artificial noise:
  - insertion, deletion, reordering

Insertion

- **Motivation**
  - Word-by-word translation always outputs a target word for every position
  - Some common words are considered as redundant ones
- **Method**
  - For each position in a sentence, insert a frequent word according from set v_{ins} to a probability distribution p_{ins}
  - Denoising network learns to delete the word when translating

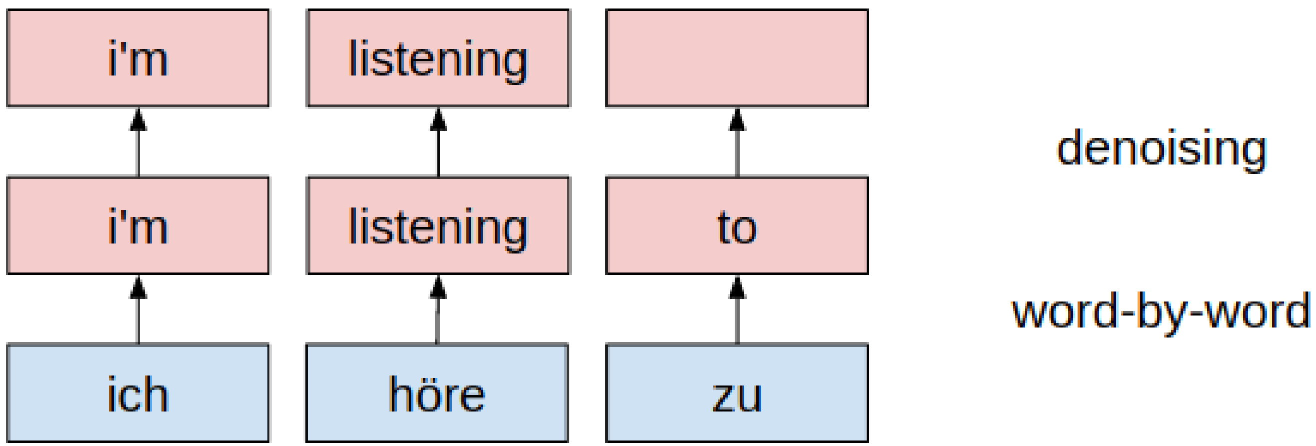

Deletion

- **Motivation**
  - **In contrary case: some words are not related to any source word**
- **Realization**
  - **For each position in a sentence, delete the word according to a probability distribution p_{del} as input**
  - **Denoising network learns to add some potential words when translating**

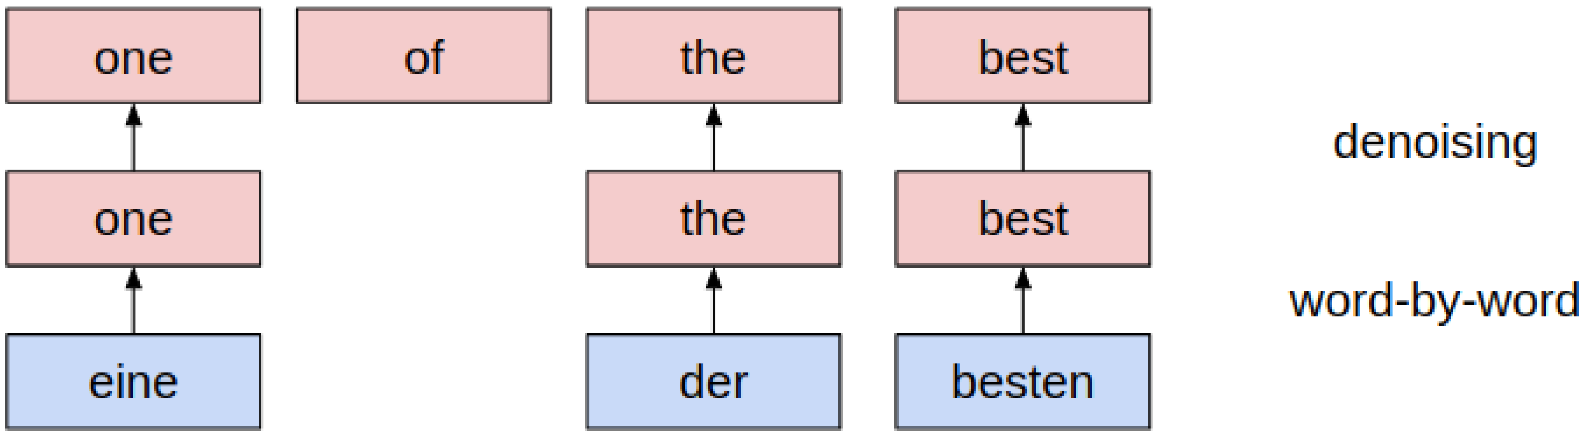

Reordering

- **Motivation**
  - **Generated words are not in a correct sequence of the target language**
- **Method**
  - **For each position of a sentence, swap the words within a limited distance d_{per} as input**
  - **Denoising network learns reordering information when translating**

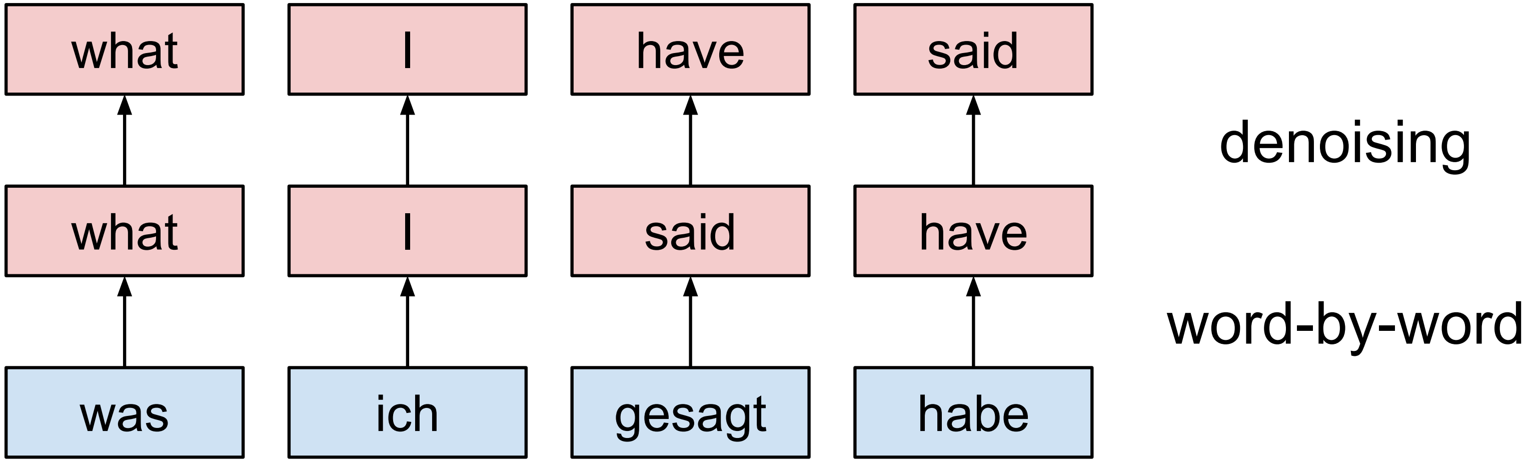

Experiment Settings

- Word embedding and LM trained on `News Crawl 2014-2017` (100M)
- BLEU evaluated on German↔English `newstest2016`
- Word accuracy evaluated on dictionaries released by Facebook
  - Dictionary built with internal translation tool
  - Each word has 1-4 word translation(s)
  - Top-1 accuracy
- Context-aware beam search
  - Lexicon candidates: 100
  - Beam width: 10

Experiments

Translation results on German↔English `newstest2016` and French↔English `newstest2014`.

System	OOV handling	de-en BLEU [%]	en-de BLEU [%]	fr-en BLEU [%]	en-fr BLEU [%]
Word-by-Word	None	11.1	6.7	10.6	7.8
+ LM (5-gram)	LM	12.9	8.9	12.7	10.0
	Copy	14.5	9.9	13.6	10.9
+ Denoising (RNN)		16.2	10.6	15.8	13.3
+ Denoising (Transformer)		17.2	11.0	16.5	13.9
[Lample & Denoyer^{+} 17]		13.3	9.6	14.3	15.1
[Artetxe & Labaka^{+} 17b]		-	-	15.6	15.1

Ablation studies

- Different sizes of training corpora
- Different vocabularies: BPE and word
- Different vocabulary sizes for cross-lingual training
- Different denoising parameters
- Phrase embedding
- Vocabulary cut-off

Different Training Corpora

Word-by-word translation from German to English

	Accuracy [%]	Bleu [%]
5M	44.9	9.7
10M	51.6	10.1
50M	59.4	10.8
100M	61.2	11.2

- Larger corpus improves the word translation accuracy
- Also improves the word-by-word translation

Different Embeddings and Training Vocabulary Size

	Vocabulary	BLEU [%]
	Merges	
BPE	20k	10.4
	50k	12.5
	100k	13.0
	Cross-lingual training	
Word	20k	14.4
	50k	14.4
	100k	14.5
	200k	14.4

- Word-by-word translation with language model
- Word embedding performs better than BPE embedding
- Embedding trained on 20k similar to 200k $\Rightarrow$ Frequent words matter

Denoising Experiments

d_{per}	p_{del}	p_{ins}	v_{ins}	BLEU [%]
2				14.7
3				14.9
5				14.9
3	0.1			15.7
	0.3			15.1
3	0.1	0.1	10	16.8
			50	17.2
			500	16.8
			5000	16.5

- Each artificial noise improves the translation performance

Phrase Embedding

Motivation

- Many phrases have a meaning that is not a simple composition of the meaning of its individual words

Phrase detection

- Phrases formed based on the unigram and bigram counts: [Mikolov & Sutskever^{+} 13]
  - Tune a good threshold value for score

$$\text{score}(e', e) = \frac{\text{count}(e', e) - \delta}{\text{count}(e') * \text{count}(e)}$$

- Process sentences with most common phrases in training corpus
  - Count the most frequent bi-gram phrases: $\text{score}(e', e) = \text{count}(e', e)$
  - Detect phrases as top frequent phrases in the training corpus

Phrase Embedding Experiments

Vocabulary			No LM BLEU [%]	With LM BLEU [%]	Denoising BLEU [%]
Word			11.2	14.5	17.2
[Mikolov & Sutskever$^+$ 13]	threshold	100	11.1	13.7	15.6
		500	11.0	13.7	16.2
		2000	10.7	14.0	16.5
Top frequent	count	50k	12.0	15.7	16.8

- Phrase embeddings helps only for WBW and $+$LM

Source and Target Vocabulary Cut-off

- Limit the vocabulary size, copy the OOV directly (mainly name entities)
- Column: source vocabulary size/ row: target vocabulary size

Word embedding vocabulary cut-off

BLEU [%]	20k	50k	100k
50k	11.1	11.3	11.2
100k	11.2	11.2	11.1
150k	10.9	10.9	-

Phrase embedding vocabulary cut-off

BLEU [%]	50k	100k	150k
50k	11.3	-	-
100k	11.9	11.9	-
150k	12.0	11.9	11.9
200k	12.0	-	-

- Vocabulary size affects the translation performance slightly

LM Supported Cross-lingual Embedding training

Motivation

- Language model help to select candidates, provide better dictionary
- Straightforward modeling, larger dictionary, different mapping types and loss functions

Model

- Training the mapping with SGD instead of Procrustes analysis
- Dictionary from the sentence translation with LM, instead of induction

$$\left.\begin{array}{c} (f_1, \quad e_1) \\ (f_2, \quad e_2) \\ \vdots \\ (f_N, \ e_N) \end{array}\right\} \Rightarrow D \qquad\qquad \mathcal{L} = \sum_{(f,e)\in D} \|Wf - e\|^2$$

Conclusions

Comprehensive results

- Context-aware beam search with LM helps the lexicon choice
- Denoising networks aimed at insertion/deletion/reordering noise works for such problems in a small range of sentences

Ablation studies

- BPE embeddings performs worse than word embeddings, especially with smaller vocabulary size.
- Word-by-word translation with cross-lingual embedding depends highly on the frequent word mappings
- Phrase embedding only helps in word-by-word translation with LM

Outlook

Goal: Improve the unsupervised learning for cross-lingual embedding

- Accordingly improves unsupervised MT performance
- Other applications: transfer learning for low-resource LM [Adams & Makarucha^{+}

Non-linear mapping between source and target

- Linear assumption may be too crude
- Stochastic gradient descent instead of SVD
- Also applies in supervised case

Thank you for your attention

Jiahui Geng

jgeng@cs.rwth-aachen.de

http://www.hltpr.rwth-aachen.de/

References

[Adams & Makarucha+ 17] O. Adams, A. Makarucha, G. Neubig, S. Bird, T. Cohn: Cross-lingual word embeddings for low-resource language modeling. In *Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics: Volume 1, Long Papers*, Vol. 1, pp. 937–947, 2017.

[Artetxe & Labaka+ 17a] M. Artetxe, G. Labaka, E. Agirre: Learning bilingual word embeddings with (almost) no bilingual data. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Vol. 1, pp. 451–462, 2017.

[Artetxe & Labaka+ 17b] M. Artetxe, G. Labaka, E. Agirre, K. Cho: Unsupervised neural machine translation. *arXiv preprint arXiv:1710.11041*, Vol., 2017.

[Artetxe & Labaka+ 18] M. Artetxe, G. Labaka, E. Agirre: A robust self-learning method for fully unsupervised cross-lingual mappings of word embeddings. *arXiv preprint arXiv:1805.06297*, Vol., 2018.

[Bojanowski & Grave+ 17] P. Bojanowski, E. Grave, A. Joulin, T. Mikolov: Enriching Word Vectors with Subword Information. *Transactions of the Association for Computational Linguistics*, Vol. 5, pp. 135–146, 2017.

[Conneau & Lample+ 17] A. Conneau, G. Lample, M. Ranzato, L. Denoyer, H. Jégou: Word translation without parallel data. *arXiv preprint arXiv:1710.04087*, Vol., 2017.

[Hoshen & Wolf 18] Y. Hoshen, L. Wolf: An Iterative Closest Point Method for Unsupervised Word Translation. *arXiv preprint arXiv:1801.06126*, Vol., 2018.

[Lample & Denoyer+ 17] G. Lample, L. Denoyer, M. Ranzato: Unsupervised Machine Translation Using Monolingual Corpora Only. *arXiv preprint arXiv:1711.00043*, Vol., 2017.

[Mikolov & Sutskever+ 13] T. Mikolov, I. Sutskever, K. Chen, G.S. Corrado, J. Dean: Distributed representations of words and phrases and their compositionality. In *Advances in neural information processing systems*, pp. 3111–3119, 2013.